

MA388 Sabermetrics: Lesson 9

Value of Plays - Run Expectancy Matrix

LTC Jim Pleuss

```
library(tidyverse)
library(Lahman)
library(knitr)
library(ggrepel)
library(broom)
```

Review

Last lesson, we discussed the following three models:

$$Wpct = \beta_0 + \beta_1 RD + \epsilon \quad (1)$$

$$Wpct = \frac{R^2}{R^2 + RA^2} + \epsilon \quad (2)$$

$$Wpct = \frac{R^k}{R^k + RA^k} + \epsilon \quad (3)$$

- Determine the value of k in the third model for the 2018 season.
- Calculate the predicted wins for each team in 2018 and plot the residuals vs. the predicted values.




Figure 1: Pythagorean predictions (2018)

Value of Plays

Using the 10-runs-per-win rule of thumb (or more precise estimates using the other models), we now have a nice way of converting runs to wins. This leads to our next obvious question – how do we score runs? Simply put, players make plays, and sometimes those plays lead to runs. It’s time to assess how many runs players and plays are worth.

Run Expectancy Matrix

* The first step is to calculate the *run expectancy matrix*. The run expectancy matrix tells us the average number of runs scored in the remainder of an inning for each state (runner/out combination) of an inning.

The state of the game consists of four digits. The first three digits indicate the location of runners on base. The last digit indicates the number of outs.

Examples:

- “000 0” - no runners on base; no outs
- “111 2” - bases loaded; two outs
- “010 1” - runner on second base; one out

Base Runners			2010-2015			1993-2009			1969-1992			1950-1968		
1B	2B	3B	0 outs	1 outs	2 outs	0 outs	1 outs	2 outs	0 outs	1 outs	2 outs	0 outs	1 outs	2 outs
—	—	—	0.481	0.254	0.098	0.547	0.293	0.113	0.477	0.252	0.094	0.476	0.256	0.098
1B	—	—	0.859	0.509	0.224	0.944	0.565	0.245	0.853	0.504	0.216	0.837	0.507	0.216
—	2B	—	1.100	0.664	0.319	1.175	0.723	0.349	1.102	0.678	0.325	1.094	0.680	0.330
1B	2B	—	1.437	0.884	0.429	1.562	0.966	0.471	1.476	0.902	0.435	1.472	0.927	0.441
—	—	3B	1.350	0.950	0.353	1.442	0.991	0.388	1.340	0.943	0.373	1.342	0.926	0.378
1B	—	3B	1.784	1.130	0.478	1.854	1.216	0.533	1.715	1.149	0.484	1.696	1.151	0.504
—	2B	3B	1.964	1.376	0.580	2.053	1.449	0.626	1.967	1.380	0.594	1.977	1.385	0.620
1B	2B	3B	2.292	1.541	0.752	2.390	1.635	0.815	2.343	1.545	0.752	2.315	1.540	0.747

Figure 2: Run Expectancy Matrices (source: <http://tangotiger.net/re24.html>)

How do we interpret the values in the matrix?

Do you see any trends over time?

We are looking at run expectancy matrices. There is another type of matrix called a [run probability matrix](#). *Explain how they differ.*

Run Expectancy Matrix from Retrosheet Data

Typically, we use Retrosheet play-by-play data to calculate the run expectancy matrix. Below, we will calculate it for 2011 using the code from our textbook.

```
# Import Retrosheet play-by-play data from our textbook site.

site = "https://raw.githubusercontent.com/maxtoki/baseball_R/"
fields <- read_csv(file = paste(site, "master/data/fields.csv", sep = ""))
retro2011 <- read_csv(file = paste(site, "master/data/all2011.csv", sep = ""),
                      col_names = pull(fields, Header),
                      na = character())
colnames(retro2011) <- tolower(colnames(retro2011))
```

This data frame contains one line for every play in the 2011 season. There are approximately 200,000 lines in this data set.

```
retro2011 |>
  select(game_id, away_team_id, inn_ct, outs_ct,
         bat_id, pit_id, event_cd) |>
  head(10) |>
  kable()
```

game_id	away_team_id	inn_ct	outs_ct	bat_id	pit_id	event_cd
ANA20110408	TOR	1	0	davir003	sante001	2
ANA20110408	TOR	1	1	nix-j001	sante001	2

game_id	away_team_id	inn_ct	outs_ct	bat_id	pit_id	event_cd
ANA20110408	TOR	1	2	bautj002	sante001	3
ANA20110408	TOR	1	0	iz-	drabk001	14
				tum001		
ANA20110408	TOR	1	0	kendh001	drabk001	3
ANA20110408	TOR	1	1	abreb001	drabk001	6
ANA20110408	TOR	1	2	abreb001	drabk001	14
ANA20110408	TOR	1	2	huntt001	drabk001	20
ANA20110408	TOR	1	2	wellv001	drabk001	2
ANA20110408	TOR	2	0	linda001	sante001	2

To understand what happened in this game, we'll need to know what all the [event codes](#) are for various baseball outcomes described by the `event_cd` field.

Code	Meaning
0	Unknown event
1	No event
2	Generic out
3	Strikeout
4	Stolen base
5	Defensive indifference
6	Caught stealing
7	Pickoff error
8	Pickoff
9	Wild pitch
10	Passed ball
11	Balk
12	Other advance
13	Foul error
14	Walk
15	Intentional walk
16	Hit by pitch
17	Interference
18	Error
19	Fielder's choice
20	Single
21	Double
22	Triple
23	Home run
24	Missing play

Figure 3: Event codes

If we further consider the `event_tx` and `pitch_seq_tx` fields, we can more comprehensively reconstruct each at-bat. Let's try to describe what happened in the first inning of this game. We'll need to know the [position numbers](#) to interpret the event text strings.




Figure 4: Baseball Position Numbers

```
retro2011 |>
  select(inn_ct, outs_ct,
         bat_id, pit_id, pitch_seq_tx, event_tx) |>
  head(10) |>
  kable()
```

inn_ct	outs_ct	bat_id	pit_id	pitch_seq_tx	event_tx
1	0	davir003	sante001	FBSX	9/F
1	1	nix- j001	sante001	X	9/F
1	2	bautj002	sante001	CBCS	K
1	0	iz- tum001	drabk001	CBBBB	W
1	0	kendh001	drabk001	BCSBS	K
1	1	abreb001	drabk001	CBB1>S	CS2(26)
1	2	abreb001	drabk001	CBB1>S.FBF	W
1	2	huntt001	drabk001	CCX	S9/L.1- H(E9/TH)(UR)(NR);B- 3
1	2	wellv001	drabk001	BX	63/G
2	0	linda001	sante001	CBBFX	3/L-

Calculate the Run Expectancy Matrix

The textbook provides the code below to generate the run expectancy matrix.

```
retro2011 <- retro2011 |>
  mutate(runs_before = away_score_ct + home_score_ct,
         half_inning = paste(game_id, inn_ct, bat_home_id),
         runs_scored =
           (bat_dest_id > 3) + (run1_dest_id > 3) +
           (run2_dest_id > 3) + (run3_dest_id > 3))

half_innings <- retro2011 |>
  group_by(half_inning) |>
  summarize(outs_inning = sum(event_outs_ct),
           runs_inning = sum(runs_scored),
           runs_start = first(runs_before),
           max_runs = runs_inning + runs_start)

retro2011 <- retro2011 |>
  inner_join(half_innings, by = "half_inning") |>
  mutate(runs_roi = max_runs - runs_before)

retro2011 <- retro2011 |>
  mutate(bases =
         paste(ifelse(base1_run_id > '', 1, 0),
              ifelse(base2_run_id > '', 1, 0),
              ifelse(base3_run_id > '', 1, 0), sep = ""),
         state = paste(bases, outs_ct))

retro2011 <- retro2011 |>
  mutate(is_runner1 =
         as.numeric(run1_dest_id == 1 | bat_dest_id == 1),
         is_runner2 =
         as.numeric(run1_dest_id == 2 | run2_dest_id == 2 |
                    bat_dest_id == 2),
         is_runner3 =
         as.numeric(run1_dest_id == 3 | run2_dest_id == 3 |
                    run3_dest_id == 3 | bat_dest_id == 3),
         new_outs = outs_ct + event_outs_ct,
         new_bases = paste(is_runner1, is_runner2, is_runner3, sep = ""),
         new_state = paste(new_bases, new_outs))
```



```

retro2011 <- retro2011 |>
  filter((state != new_state) | (runs_scored > 0))

retro2011_complete <- retro2011 |>
  filter(outs_inning == 3)

erm_2011 <- retro2011_complete |>
  group_by(bases, outs_ct) |> # same as grouping by state
  summarize(mean_run_value = mean(runs_roi))

erm_2011_wide <- erm_2011 |>
  pivot_wider(
    names_from = outs_ct,
    values_from = mean_run_value,
    names_prefix = "Outs="
  )

write_csv(erm_2011_wide, "erm_2011_wide.csv")

erm_2011_wide |> kable()

```

bases	Outs=0	Outs=1	Outs=2
000	0.4711649	0.2546956	0.0971845
001	1.4543568	0.9374359	0.3173913
010	1.0582804	0.6501976	0.3091673
011	1.9304897	1.3388641	0.5407407
100	0.8350992	0.4960492	0.2179536
101	1.7526998	1.1496169	0.4882597
110	1.4144549	0.8739176	0.4222569
111	2.1718266	1.4745146	0.7610094

Now we use this expected run table to determine the value of every event.

```

retro2011 <- retro2011 |>
  left_join(erm_2011, join_by("bases", "outs_ct")) |>
  rename(rv_start = mean_run_value) |>
  left_join(
    erm_2011,
    join_by(new_bases == bases, new_outs == outs_ct)
  ) |>
  rename(rv_end = mean_run_value) |>
  replace_na(list(rv_end = 0)) |>
  mutate(run_value = rv_end - rv_start + runs_scored)

```

Note the code above also adds some very useful columns to the play-by-play data related to the state and run expectancy at the start and end of each play.

```
retro2011 |>
  select(bat_id, event_cd, state, rv_start,
         new_state, rv_end, runs_scored, run_value) |>
  head(10) |>
  kable(digits = 3)
```

bat_id	event_cd	state	rv_start	new_state	rv_end	runs_scored	run_value
davir003	2	000 0	0.471	000 1	0.255	0	-0.216
nix- j001	2	000 1	0.255	000 2	0.097	0	-0.158
bautj002	3	000 2	0.097	000 3	0.000	0	-0.097
iz- tum001	14	000 0	0.471	100 0	0.835	0	0.364
kendh001	3	100 0	0.835	100 1	0.496	0	-0.339
abreb001	6	100 1	0.496	000 2	0.097	0	-0.399
abreb001	14	000 2	0.097	100 2	0.218	0	0.121
huntt001	20	100 2	0.218	001 2	0.317	1	1.099
wellv001	2	001 2	0.317	001 3	0.000	0	-0.317
linda001	2	000 0	0.471	000 1	0.255	0	-0.216

```
library(baseballr) #the appendix does not tell you that you need this.

retro_data<-retrosheet_data(years=2022)
retro2022 <- retro_data |>
  pluck("2022") |>
  pluck("events")
roster2022 <- retro_data |>
  pluck("2022") |>
  pluck("rosters")
roster2022_clean<-roster2022 |>
  select(player_id,last_name,first_name) |>
  rename(bat_id=player_id) |>
```

```

distinct()

stats<-retro2022 |> #I split the code chunk so I don't have to do the slow step every time
select(bat_id,event_cd,bat_home_id) |> #grab only the columns I want
filter(event_cd>=20|event_cd==2|event_cd==3) |> #events >= 20 are hits, 2+3 are out/strikeouts
group_by(bat_id,bat_home_id) |>
mutate(
hit=if_else(event_cd>=20,1,0), #if play event was hit make it 1 if not 0
ab=1, #all of these are ABs, making them 1 to sum later
home_away=if_else(bat_home_id==1,"Home","Away"))

player_stats<-stats |>
group_by(bat_id,home_away) |>
summarize(AB=sum(ab),H=sum(hit),AVG=H/AB) |>
filter(AB>=100) #filter out players with <100 ABs

player_home_v_away<-player_stats |>
select(bat_id,home_away,AVG) |>
pivot_wider(id_cols = bat_id,
names_from = home_away,
values_from=AVG,
)

roster2022 |> count(player_id) |> arrange(-n)

```

A tibble: 1,495 x 2

	player_id	n
	<chr>	<int>
1	chany001	4
2	fordm002	4
3	abrea001	3
4	banda001	3
5	canor001	3
6	chavj001	3
7	duggs001	3
8	fairs001	3
9	gonzc002	3
10	keucd001	3

i 1,485 more rows

```

player_home_v_away<-player_home_v_away |>
left_join(roster2022_clean,by="bat_id")
player_home_v_away |>
mutate(Difference=abs(Home-Away)) |>
arrange(Difference) |>

```

```
head(10) |>
select(last_name,first_name,Difference)
```

```
# A tibble: 10 x 4
# Groups:   bat_id [10]
  bat_id last_name first_name Difference
  <chr>   <chr>    <chr>      <dbl>
1 darnt001 d'Arnaud  Travis      0
2 vierm001 Vierling  Matt      0.000149
3 andre001 Andrus   Elvis      0.000569
4 wendj002 Wendle   Joey       0.000605
5 santc002 Santana  Carlos     0.00110
6 merrw001 Merrifield Whit      0.00151
7 heimj001 Heim     Jonah      0.00156
8 smitp002 Smith    Pavin      0.00192
9 estrt001 Estrada  Thairo     0.00196
10 rizza001 Rizzo    Anthony    0.00208
```