

MA388 Sabermetrics: Lesson 16

Catcher Framing Ability Random Effects

LTC Jim Pleuss

Last class, we were primarily comparing called strikes for Buster Posey and Yadier Molina. Today we'll look at more catchers from June 2019.

```
library(tidyverse)
library(knitr)
library(broom)
library(mgcv)

# Retrieve pitch level data from June 2019. (See previous lesson.)
pitches <- readRDS("statcast_june_2019.rds")

# Add catcher's name to the pitch data from the MLB master list.
mlbIDs <- baseballr::chadwick_player_lu() |>
  mutate(
    mlb_name = paste(name_first, name_last),
    mlb_id = key_mlbam
  )

pitches <- pitches |>
  left_join(select(mlbIDs, mlb_name, mlb_id),
            by = c("fielder_2" = "mlb_id")) |>
  rename(catcher_name = mlb_name)
#
# pitches_taken_subset <- pitches |>
#   filter(description %in% c("ball", "called_strike"))
```

Random Effects Models

Thus far, we investigated the difference between Molina and Posey on called strikes. In doing so, we did the following:

1. Compared their called strike proportions (*why isn't this that useful?*)

2. Estimated the log odds ratio for a called strike after adjusting for pitch location (*why is this better?*)

But, what we really want to know is not if Molina is better than Posey, but if catchers, in general, have a meaningful effect on called strikes.

Based on what you've learned so what far, you would probably fit a *fixed effects* model like this:

$$\log\left(\frac{\pi_i}{1-\pi_i}\right) = \beta_0 + \beta_1 \text{catcher}_{1,i} + \dots + \beta_{m-1} \text{catcher}_{m-1,i} + f(\text{plate_x}_i, \text{plate_z}_i)$$

where $\text{catcher}_{j,i}$ is an indicator of whether the catcher for pitch i is catcher j .

How many parameters are there in the model above?

Let's fit the model for all catchers who caught at least 1000 pitches:

```
catcher_list <- pitches |>
  group_by(catcher_name) |>
  summarize(n = n()) |>
  filter(n >= 1000) |>
  pull(catcher_name)

pitches_taken <- pitches |>
  filter(catcher_name %in% catcher_list,
         description %in% c("called_strike", "ball")) |>
  mutate(called_strike=description=='called_strike')

strike_mod_all <- gam(called_strike ~ s(plate_x, plate_z) + catcher_name,
                     family = "binomial", data = pitches_taken)

summary(strike_mod_all)
```

Family: binomial

Link function: logit

Formula:

called_strike ~ s(plate_x, plate_z) + catcher_name

Parametric coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.622732	1.142689	-5.796	6.8e-	
09 ***					
catcher_nameAustin Hedges	0.504288	0.202702	2.488	0.01285	*
catcher_nameBobby Wilson	0.029825	0.245906	0.121	0.90346	
catcher_nameBrian McCann	0.056919	0.213540	0.267	0.78982	
catcher_nameBryan Holaday	-0.256764	0.220585	-1.164	0.24442	
catcher_nameBuster Posey	-0.257745	0.215861	-1.194	0.23247	
catcher_nameCam Gallagher	0.056798	0.241117	0.236	0.81377	
catcher_nameCarson Kelly	0.144979	0.199848	0.725	0.46818	
catcher_nameChance Sisco	-0.473557	0.208069	-2.276	0.02285	*
catcher_nameChris Iannetta	-0.390696	0.218313	-1.790	0.07352	.
catcher_nameChristian Vázquez	0.443846	0.196247	2.262	0.02372	*
catcher_nameCurt Casali	-0.190821	0.205608	-0.928	0.35337	
catcher_nameDanny Jansen	0.016908	0.198817	0.085	0.93223	
catcher_nameDustin Garneau	-0.037663	0.242865	-0.155	0.87676	
catcher_nameEliás Díaz	-0.571133	0.191465	-2.983	0.00285	**
catcher_nameGary Sánchez	0.085060	0.192510	0.442	0.65860	
catcher_nameGrayson Greiner	-0.337244	0.259916	-1.298	0.19446	
catcher_nameJ. T. Realmuto	-0.052478	0.192935	-0.272	0.78562	
catcher_nameJames McCann	-0.548860	0.200074	-2.743	0.00608	**
catcher_nameJason Castro	-0.059725	0.208673	-0.286	0.77471	
catcher_nameJeff Mathis	-0.209227	0.204925	-1.021	0.30726	
catcher_nameJohn Hicks	-0.360805	0.218333	-1.653	0.09842	.
catcher_nameJonathan Lucroy	-0.434742	0.197432	-2.202	0.02767	*
catcher_nameJorge Alfaro	0.302585	0.214719	1.409	0.15877	
catcher_nameJosh Phegley	-0.418253	0.191059	-2.189	0.02859	*
catcher_nameKevin Plawecki	0.555562	0.250585	2.217	0.02662	*
catcher_nameKurt Suzuki	-0.230635	0.212513	-1.085	0.27780	
catcher_nameLuke Maile	-0.340531	0.219758	-1.550	0.12124	
catcher_nameManny Piña	0.532831	0.242093	2.201	0.02774	*
catcher_nameMartín Maldonado	-0.099924	0.196022	-0.510	0.61022	
catcher_nameMatt Wieters	-0.330689	0.229847	-1.439	0.15023	
catcher_nameMike Zunino	-0.040333	0.206641	-0.195	0.84525	
catcher_nameMitch Garver	-0.025173	0.214782	-0.117	0.90670	
catcher_nameOmar Narváez	-0.390007	0.195070	-1.999	0.04557	*
catcher_namePedro Severino	-0.519941	0.201742	-2.577	0.00996	**
catcher_nameRoberto Pérez	0.101593	0.198618	0.512	0.60900	
catcher_nameRobinson Chirinos	-0.317570	0.191141	-1.661	0.09662	.
catcher_nameRussell Martin	0.224721	0.218288	1.029	0.30326	
catcher_nameSandy León	0.090762	0.222809	0.407	0.68375	
catcher_nameStephen Vogt	-0.200628	0.217149	-0.924	0.35553	
catcher_nameTim Federowicz	0.055533	0.220292	0.252	0.80097	

```

catcher_nameTom Murphy      0.251783    0.211580    1.190    0.23404
catcher_nameTomás Nido      0.084321    0.228749    0.369    0.71241
catcher_nameTony Wolters    -0.300176    0.196645   -1.526    0.12689
catcher_nameTravis d'Arnaud -0.117914    0.214300   -0.550    0.58216
catcher_nameTucker Barnhart -0.002438    0.222311   -0.011    0.99125
catcher_nameTyler Flowers   0.436709    0.213844    2.042    0.04113 *
catcher_nameVíctor Caratini  0.111638    0.251220    0.444    0.65676
catcher_nameWillson Contreras -0.077984    0.191878   -0.406    0.68443
catcher_nameWilson Ramos    -0.055399    0.196864   -0.281    0.77840
catcher_nameYadier Molina   -0.106296    0.199854   -0.532    0.59482
catcher_nameYan Gomes       -0.126179    0.209642   -0.602    0.54725
catcher_nameYasmani Grandal  0.156120    0.193919    0.805    0.42077
---

```

```

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Approximate significance of smooth terms:

```

              edf Ref.df Chi.sq p-value
s(plate_x,plate_z) 27.26  28.01   8062  <2e-16 ***
---

```

```

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

R-sq.(adj) =  0.757   Deviance explained = 72.9%
UBRE = -0.65503   Scale est. = 1           n = 54411

```

What should we conclude from this model?

What are some limitations of a *fixed effects* approach for this question?

Random effects model the individual catcher effects as normally distributed with mean 0 and variance σ^2 . There is only one parameter to fit. As a bonus, it has a nice interpretation...the larger the variance, the larger the effect of the catcher (*why?*). Here is the model:

$$\log\left(\frac{\pi_i}{1-\pi_i}\right) = \beta_0 + \text{catcher}_j + f(\text{plate_x}_i, \text{plate_z}_i)$$

where catcher_j is normally distributed with mean 0 and variance σ^2 .

How many parameters are in this model?

Why is this an improvement?

We'll fit the model using the `lme4` package. Note, we'll do this slightly differently by first fitting the location model and then using the predictions from that model in a second model with the random effect.

```
model_location <- gam(called_strike ~ s(plate_x, plate_z),
                      family = "binomial",
                      data = pitches_taken)

pitches_taken <- model_location |>
  augment(type.predict = "response",
          newdata = pitches_taken) |>
  rename(strike_prob = .fitted)

# Now fit a random effects model adjusting for pitch location.
library(lme4)
model_random_effects_catcher <- glmer(called_strike ~ strike_prob + (1|catcher_name),
                                       family = "binomial",
                                       data = pitches_taken)
summary(model_random_effects_catcher)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: binomial ( logit )
Formula: called_strike ~ strike_prob + (1 | catcher_name)
Data: pitches_taken
```

AIC	BIC	logLik	-2*log(L)	df.resid
20889.2	20916.0	-10441.6	20883.2	54408

Scaled residuals:

Min	1Q	Median	3Q	Max
-6.9648	-0.1583	-0.1373	0.1717	8.1197

Random effects:

Groups	Name	Variance	Std.Dev.
catcher_name	(Intercept)	0.05715	0.2391

Number of obs: 54411, groups: catcher_name, 53

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.89584	0.04674	-83.36	<2e-16 ***
strike_prob	7.48119	0.05805	128.88	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Correlation of Fixed Effects:

(Intr)
strike_prob -0.572

```
# Get random effects for catchers.  
catcher_effects_adj <- model_random_effects_catcher |>  
  ranef() |>  
  as_tibble() |>  
  transmute(id = levels(grp), effect = condval) |>  
  arrange(desc(effect))  
  
catcher_effects_adj |> head(5)
```

```
# A tibble: 5 x 2  
  id          effect  
  <chr>        <dbl>  
1 James McCann 0.432  
2 Tony Wolters 0.366  
3 Wilson Ramos 0.364  
4 Jason Castro 0.359  
5 Jorge Alfaro 0.343
```

```
catcher_effects_adj |>  
  ggplot(aes(x = effect)) +  
  geom_histogram()
```



Interesting - I guess? We're now faced with the same question we always have in statistics: I see a statistic, a non-zero difference, a non-zero variance, etc., but is it significant? What are some ways we can go about determining significance?

Let's use a nested model approach and add a couple more explanatory variables. Right now we've only let the catcher and the pitch location try to explain called strike determinations. If we introduce other variables that provide the same or complementary information, we might find the importance of catcher in the model to decrease. Sounds like ANOVA, right?

```
model_random_effects_trio <- glmer(called_strike ~ strike_prob +
                                   (1|pitcher) + (1|batter) + (1|catcher_name),
                                   family = "binomial",
                                   data = pitches_taken)
summary(model_random_effects_trio)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: binomial ( logit )
Formula: called_strike ~ strike_prob + (1 | pitcher) + (1 | batter) +
(1 | catcher_name)
```

Data: pitches_taken

AIC	BIC	logLik	-2*log(L)	df.resid
20829.2	20873.7	-10409.6	20819.2	54406

Scaled residuals:

Min	1Q	Median	3Q	Max
-8.0613	-0.1588	-0.1307	0.1659	8.8208

Random effects:

Groups	Name	Variance	Std.Dev.
batter	(Intercept)	0.05828	0.2414
pitcher	(Intercept)	0.07996	0.2828
catcher_name	(Intercept)	0.04701	0.2168

Number of obs: 54411, groups: batter, 600; pitcher, 531; catcher_name, 53

Fixed effects:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.98030	0.04928	-80.77	<2e-16 ***
strike_prob	7.63002	0.05965	127.92	<2e-16 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)
strike_prob	-0.560

```
# Get random effects for pitchers, catchers, and batters.
player_effects_adj <- model_random_effects_trio |>
  ranef() |>
  as_tibble() |>
  mutate(id = grp, effect = condval) |>
  arrange(desc(effect))

player_effects_adj |>
  ggplot(aes(x = effect)) +
  geom_histogram() +
  facet_wrap(~grpvar, scales = "free_y")
```

Let's talk course project!