

MA388 Sabermetrics: Lesson 3

Introduction to R - Part II

```
library(Lahman)
library(tidyverse)
library(knitr)
```

Review

Key Concepts

- five tidyverse verbs
- relational databases

Review Questions

- Using the **Teams** data frame, calculate average wins (W) per season for each MLB team from 2000-present. Order your results from the best teams to the worst.

teamID	W	G	Wpct
NYA	2286	3946	0.579
LAN	2239	3948	0.567
SLN	2179	3945	0.552
ATL	2152	3945	0.546
BOS	2145	3947	0.543
ANA	425	810	0.525
SFN	2052	3946	0.520
CLE	2046	3945	0.519
OAK	2030	3946	0.514
HOU	2029	3947	0.514
LAA	1609	3138	0.513
PHI	2019	3947	0.512

teamID	W	G	Wpct
MIN	1984	3949	0.502
TOR	1977	3947	0.501
NYN	1976	3946	0.501
SEA	1975	3948	0.500
CHN	1969	3947	0.499
TBA	1959	3946	0.496
MIL	1960	3949	0.496
FLO	963	1942	0.496
TEX	1937	3949	0.491
ARI	1922	3948	0.487
CHA	1911	3948	0.484
WAS	1510	3136	0.482
SDN	1888	3949	0.478
CIN	1861	3950	0.471
DET	1841	3943	0.467
COL	1809	3949	0.458
BAL	1805	3948	0.457
MON	368	810	0.454
PIT	1767	3944	0.448
KCA	1737	3948	0.440
MIA	868	2002	0.434

- Using the `Teams` data frame, calculate how many World Series titles (WSWin) each team had during this period and add this information to your table in the previous question.

Table 2: Win Percentage and World Series Titles by Team (2000-present).

teamID	W	G	Wpct	WSWins
NYA	2286	3946	0.579	2
LAN	2239	3948	0.567	2
SLN	2179	3945	0.552	2
ATL	2152	3945	0.546	1
BOS	2145	3947	0.543	4
ANA	425	810	0.525	1
SFN	2052	3946	0.520	3
CLE	2046	3945	0.519	0
OAK	2030	3946	0.514	0
HOU	2029	3947	0.514	2
LAA	1609	3138	0.513	0
PHI	2019	3947	0.512	1
MIN	1984	3949	0.502	0

teamID	W	G	Wpct	WSWins
TOR	1977	3947	0.501	0
NYN	1976	3946	0.501	0
SEA	1975	3948	0.500	0
CHN	1969	3947	0.499	1
TBA	1959	3946	0.496	0
MIL	1960	3949	0.496	0
FLO	963	1942	0.496	1
TEX	1937	3949	0.491	1
ARI	1922	3948	0.487	1
CHA	1911	3948	0.484	1
WAS	1510	3136	0.482	1
SDN	1888	3949	0.478	0
CIN	1861	3950	0.471	0
DET	1841	3943	0.467	0
COL	1809	3949	0.458	0
BAL	1805	3948	0.457	0
MON	368	810	0.454	0
PIT	1767	3944	0.448	0
KCA	1737	3948	0.440	1
MIA	868	2002	0.434	0

Introduction to R (Day 2)

Split, Apply, and Combine Data

It's important in data science to be able to write procedures and apply them over a data set. The Marchi text refers to this as “splitting, applying, and combining data.” In other words, we split a data frame into pieces, apply a procedure or function to each piece, and then combine the results into a new data frame.

Ways to “split, apply, and combine”...

- `group_by()` and `summarize()` functions
- `split()` and `map_df()` functions

In general, these coding tasks involve four steps:

1. Write a function to perform the task on one split of the data.
2. Pull the individual values of the variable you are splitting on.
3. Split the data on a variable.
4. Apply the function to each split of the data.

Example

Find the team with the most wins for each season and determine whether the team won the World Series.

Four steps:

1. Write a function that takes a data frame and returns the team with the most wins and whether they won the World Series.
2. Pull the `yearID`s you will split on.
3. Split the data frame on `year`.
4. Apply the function to each split.

```
# Step 1 - Write a function.
mostWins <- function(data){
  data |>
    arrange(-W) |>
    select(teamID, W, WSWin) |>
    head(1)
}

# Step 2 - Pull the years you will split on.
year <- Teams |>
  filter(yearID >= 2000) |>
  group_by(yearID) |>
  group_keys() |>
  pull(yearID)

# Steps 3 and 4 - Split and apply.
winLeaders <- Teams |>
  filter(yearID >= 2000) |>
  split(year) |>
  map_df(mostWins, .id = "yearID")

winLeaders |>
  kable(caption = "Major League win leaders by season and whether they won the World Series.")
```

Table 3: Major League win leaders by season and whether they won the World Series.

yearID	teamID	W	WSWin
2000	SFN	97	N

yearID	teamID	W	WSWin
2001	SEA	116	N
2002	NYA	103	N
2003	ATL	101	N
2004	SLN	105	N
2005	SLN	100	N
2006	NYN	97	N
2007	BOS	96	Y
2008	LAA	100	N
2009	NYA	103	Y
2010	PHI	97	N
2011	PHI	102	N
2012	WAS	98	N
2013	BOS	97	Y
2014	LAA	98	N
2015	SLN	100	N
2016	CHN	103	Y
2017	LAN	104	N
2018	BOS	108	Y
2019	HOU	107	N
2020	LAN	43	Y
2021	SFN	107	N
2022	LAN	111	N
2023	ATL	104	N
2024	LAN	98	Y

Note: Functions are not executed until they are called.

Your Turn

Find the team with the most home runs for each season from 2000-present. Include the team's full name in your table instead of just the teamID.

Table 4: Major League win leaders by season and whether they won the World Series.

yearID	name	HR
2000	Houston Astros	249
2001	Texas Rangers	246
2002	Texas Rangers	230
2003	Texas Rangers	239
2004	Chicago White Sox	242
2005	Texas Rangers	260

yearID	name	HR
2006	Chicago White Sox	236
2007	Milwaukee Brewers	231
2008	Chicago White Sox	235
2009	New York Yankees	244
2010	Toronto Blue Jays	257
2011	New York Yankees	222
2012	New York Yankees	245
2013	Baltimore Orioles	212
2014	Baltimore Orioles	211
2015	Toronto Blue Jays	232
2016	Baltimore Orioles	253
2017	New York Yankees	241
2018	New York Yankees	267
2019	Minnesota Twins	307
2020	Los Angeles Dodgers	118
2021	Toronto Blue Jays	262
2022	New York Yankees	254
2023	Atlanta Braves	307
2024	New York Yankees	237

What's the difference between `select()` and `pull()`? One key difference is they return different data types.

```
HRLeaders |>
  select(yearID)
```

```
   yearID
1    2000
2    2001
3    2002
4    2003
5    2004
6    2005
7    2006
8    2007
9    2008
10   2009
11   2010
12   2011
13   2012
14   2013
15   2014
16   2015
```

```
17 2016
18 2017
19 2018
20 2019
21 2020
22 2021
23 2022
24 2023
25 2024
```

```
HRLeaders |>
  pull(yearID)
```

```
[1] "2000" "2001" "2002" "2003" "2004" "2005" "2006" "2007" "2008" "2009"
[11] "2010" "2011" "2012" "2013" "2014" "2015" "2016" "2017" "2018" "2019"
[21] "2020" "2021" "2022" "2023" "2024"
```